

## U23AI118

### Lab 2 : Analyzing NYC Dataset

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

df = pd.read_csv('yellow_tripdata_sample.csv')
print("Shape = ", df.shape)
print("Columns:", df.columns.tolist()) # tolist gives proper structure
print("Initial row", df.head())

Shape = (995, 19)
Columns: ['VendorID', 'tpep_pickup_datetime', 'tpep_dropoff_datetime', 'passenger_count', 'trip_distance', 'RatecodeID', 'store_and_fwd_flag', 'PULocationID', 'DOLocationID', 'payment_type', 'fare_amount', 'extra', 'mta_tax', 'tip_amount', 'tolls_amount', 'improvement_surcharge', 'total_amount', 'congestion_surcharge', 'Airport_fee']
Initial row
  VendorID tpep_pickup_datetime tpep_dropoff_datetime passenger_count \
0         2 2024-01-01 00:57:55 2024-01-01 01:17:43             1.0
1         1 2024-01-01 00:03:00 2024-01-01 00:09:36             1.0
2         1 2024-01-01 00:17:06 2024-01-01 00:35:01             1.0
3         1 2024-01-01 00:36:38 2024-01-01 00:44:56             1.0
4         1 2024-01-01 00:46:51 2024-01-01 00:52:57             1.0

  trip_distance RatecodeID store_and_fwd_flag PULocationID DOLocationID \
0           1.72          1.0                N           186           79
1           1.80          1.0                N           140          236
2           4.70          1.0                N           236           79
3           1.40          1.0                N            79          211
4           0.80          1.0                N           211          148

  payment_type fare_amount extra mta_tax tip_amount tolls_amount \
0             2          17.7    1.0    0.5         0.00         0.0
1             1          10.0    3.5    0.5         3.75         0.0
2             1          23.3    3.5    0.5         3.00         0.0
3             1          10.0    3.5    0.5         2.00         0.0
4             1           7.9    3.5    0.5         3.20         0.0

  improvement_surcharge total_amount congestion_surcharge Airport_fee
0                   1.0         22.70                 2.5         0.0
1                   1.0         18.75                 2.5         0.0
2                   1.0         31.30                 2.5         0.0
3                   1.0         17.00                 2.5         0.0
4                   1.0         16.10                 2.5         0.0

# Analysing before Data Cleaning
print("Info :",df.info())
print("Missing Values :")
```

```

print(df.isnull().sum())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 995 entries, 0 to 994
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   VendorID                             995 non-null    int64
1   tpep_pickup_datetime                 995 non-null    object
2   tpep_dropoff_datetime                995 non-null    object
3   passenger_count                      995 non-null    float64
4   trip_distance                       995 non-null    float64
5   RatecodeID                          995 non-null    float64
6   store_and_fwd_flag                  995 non-null    object
7   PULocationID                        995 non-null    int64
8   DOLocationID                        995 non-null    int64
9   payment_type                        995 non-null    int64
10  fare_amount                         995 non-null    float64
11  extra                              995 non-null    float64
12  mta_tax                             995 non-null    float64
13  tip_amount                          995 non-null    float64
14  tolls_amount                        995 non-null    float64
15  improvement_surcharge               995 non-null    float64
16  total_amount                        995 non-null    float64
17  congestion_surcharge               995 non-null    float64
18  Airport_fee                         995 non-null    float64
dtypes: float64(12), int64(4), object(3)
memory usage: 147.8+ KB
Info : None
Missing Values :
VendorID                0
tpep_pickup_datetime    0
tpep_dropoff_datetime   0
passenger_count         0
trip_distance           0
RatecodeID              0
store_and_fwd_flag      0
PULocationID            0
DOLocationID            0
payment_type            0
fare_amount             0
extra                   0
mta_tax                 0
tip_amount              0
tolls_amount            0
improvement_surcharge   0

```

```

total_amount            0
congestion_surcharge    0
Airport_fee             0
dtype: int64

df_clean = df.copy()

# Filter out invalid values
df_clean = df_clean[
    (df_clean['passenger_count'] > 0) &
    (df_clean['passenger_count'] <= 8) &
    (df_clean['trip_distance'] > 0) &
    (df_clean['trip_distance'] <= 100) &
    (df_clean['fare_amount'] > 0) &
    (df_clean['fare_amount'] <= 500) &
    (df_clean['total_amount'] > 0) &
    (df_clean['total_amount'] <= 500) &
    (df_clean['tip_amount'] >= 0) &
    (df_clean['tip_amount'] <= 100) &
    (df_clean['extra'] >= 0)
]

print(f"Original dataset size: {len(df)}")
print(f"Cleaned dataset size: {len(df_clean)}")
print(f"Removed {len(df) - len(df_clean)} rows ({((len(df) - len(df_clean))/len(df)*100):.2f}%")

Original dataset size: 995
Cleaned dataset size: 968
Removed 27 rows (2.71%)

df.describe()

```

	VendorID	passenger_count	trip_distance	RatecodeID	PULocationID	\
count	995.000000	995.000000	995.000000	995.000000	995.000000	
mean	1.703518	1.581910	2.890472	1.138693	165.323618	
std	0.456936	0.911625	3.125268	3.126260	67.808656	
min	1.000000	0.000000	0.000000	1.000000	4.000000	
25%	1.000000	1.000000	1.125000	1.000000	114.000000	
50%	2.000000	1.000000	1.960000	1.000000	161.000000	
75%	2.000000	2.000000	3.400000	1.000000	236.000000	
max	2.000000	6.000000	23.900000	99.000000	265.000000	

	DOLocationID	payment_type	fare_amount	extra	mta_tax	\
count	995.000000	995.000000	995.000000	995.000000	995.000000	
mean	163.204020	1.228141	18.278593	1.741457	0.491960	
std	73.059288	0.486448	14.916422	1.255742	0.077276	
min	4.000000	1.000000	-47.800000	-1.000000	-0.500000	
25%	107.000000	1.000000	9.300000	1.000000	0.500000	

50%	162.000000	1.000000	14.200000	1.000000	0.500000
75%	236.000000	1.000000	22.600000	3.500000	0.500000
max	265.000000	4.000000	180.000000	7.750000	0.500000

	tip_amount	tolls_amount	improvement_surcharge	total_amount	\
count	995.000000	995.000000	995.000000	995.000000	
mean	3.483146	0.197357	0.988945	26.847538	
std	4.195507	1.180907	0.144929	18.092018	
min	0.000000	0.000000	-1.000000	-52.800000	
25%	0.000000	0.000000	1.000000	16.320000	
50%	3.000000	0.000000	1.000000	22.200000	
75%	4.820000	0.000000	1.000000	31.770000	
max	80.000000	12.750000	1.000000	217.200000	

	congestion_surcharge	Airport_fee
count	995.000000	995.000000
mean	2.326633	0.052764
std	0.673846	0.309523
min	-2.500000	-1.750000
25%	2.500000	0.000000
50%	2.500000	0.000000
75%	2.500000	0.000000
max	2.500000	1.750000

```
numerical_cols = ['passenger_count', 'trip_distance', 'fare_amount', 'total_amount', 'tip_amount']
def descriptive_stats(data, column):    #Description of data function
    stats_dict = {
        'mean': data[column].mean(),
        'median': data[column].median(),
        'mode': data[column].mode().iloc[0] if not data[column].mode().empty else np.nan,
        'min': data[column].min(),
        'max': data[column].max(),
        'std': data[column].std(),
        'variance': data[column].var(),
        'skew': data[column].skew(),
        'kurtosis': data[column].kurtosis()
    }
    return stats_dict
stats_summary = {}
for col in numerical_cols:
    stats_summary[col] = descriptive_stats(df_clean, col)

stats_df = pd.DataFrame(stats_summary).round(4)    # pd.DataFrame() - Converting into tabular format

stats_df
```

passenger_count	trip_distance	fare_amount	total_amount	\
-----------------	---------------	-------------	--------------	---

mean	1.6033	2.9154	18.4021	27.0301
median	1.0000	1.9800	14.2000	22.2000
mode	1.0000	0.9000	6.5000	16.3200
min	1.0000	0.0100	3.0000	7.6000
max	6.0000	23.9000	125.5000	165.6400
std	0.9037	3.1103	13.5733	16.4536
variance	0.8166	9.6739	184.2334	270.7200
skew	1.7247	3.3208	2.5125	2.6814
kurtosis	2.8847	14.2673	10.2395	11.5664

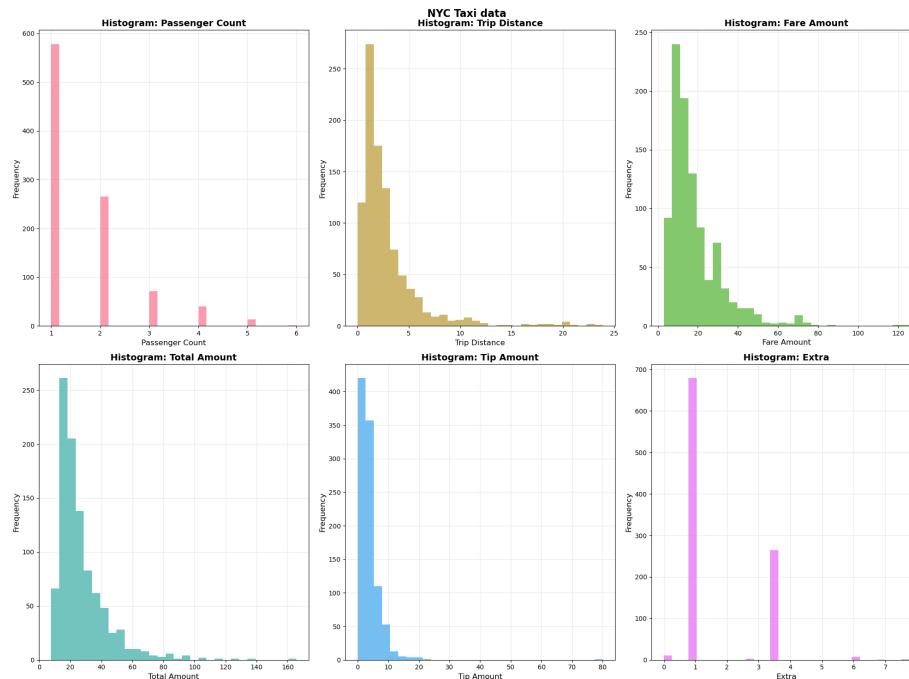
	tip_amount	extra
mean	3.4819	1.7371
median	3.0000	1.0000
mode	0.0000	1.0000
min	0.0000	0.0000
max	80.0000	7.7500
std	4.0774	1.2323
variance	16.6254	1.5186
skew	7.6468	1.2739
kurtosis	129.1305	0.9350

```
fig, axes = plt.subplots(2, 3, figsize = (20,15))
fig.suptitle('NYC Taxi data', fontsize = 16, fontweight = 'bold')

axes = axes.flatten()

for i, col in enumerate(numerical_cols):
    axes[i].hist(df_clean[col], bins =30 , alpha = 0.7, color=sns.color_palette("husl", len(numerical_cols)))
    axes[i].set_title(f'Histogram: {col.replace("_", " ").title()}', fontsize = 14, fontweight = 'bold')
    axes[i].set_xlabel(col.replace("_", " ").title(), fontsize = 12)
    axes[i].set_ylabel('Frequency', fontsize =12)
    axes[i].grid(True, alpha = 0.3)

plt.tight_layout()
plt.show()
```



All graphs are right skewed

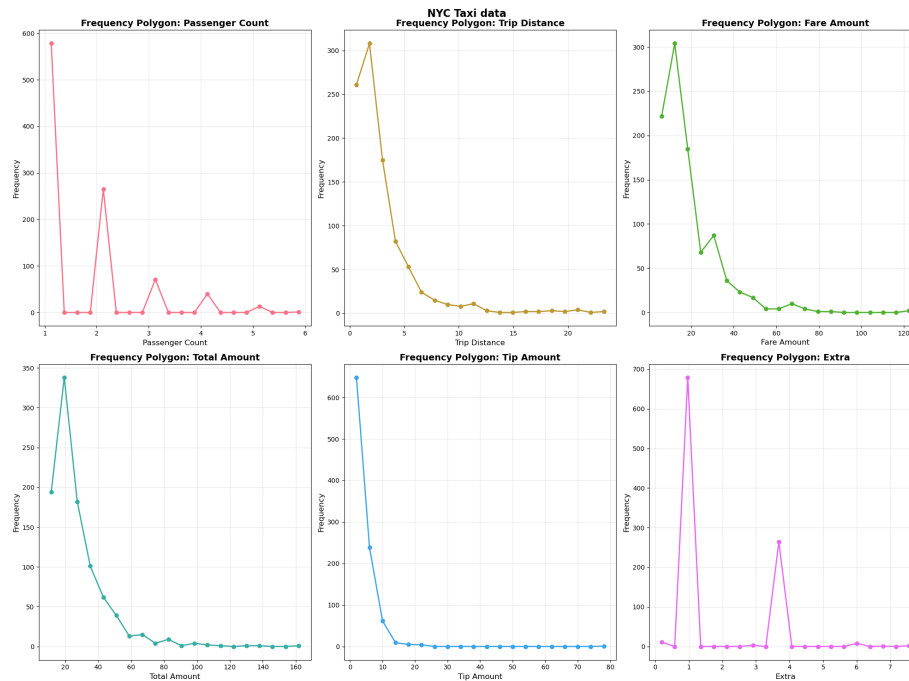
- 1) Single passenger travels more
- 2) Most people travel within range of 0-5 km
- 3) Generally fares are below \$40
- 4) Total amount comes between \$0 - \$60
- 5) people generally tips below \$ 10
- 6) Extras are either 1 or 4

```
fig, axes = plt.subplots(2, 3, figsize = (20,15))
fig.suptitle('NYC Taxi data', fontsize = 16, fontweight = 'bold')
```

```
axes = axes.flatten()
```

```
for i, col in enumerate(numerical_cols):
    hist_data, bin_edges = np.histogram(df_clean[col], bins=20)
    bin_centers = (bin_edges[:-1] + bin_edges[1:]) / 2
    axes[i].plot(bin_centers, hist_data, marker='o', linestyle='-',
                 linewidth=2, markersize=6,
                 color=sns.color_palette("husl", len(numerical_cols))[i])
    axes[i].set_title(f'Frequency Polygon: {col.replace("_", " ").title()}', fontsize = 14, fontweight = 'bold')
    axes[i].set_xlabel(col.replace("_", " ").title(), fontsize = 12)
    axes[i].set_ylabel('Frequency', fontsize = 12)
    axes[i].grid(True, alpha=0.3)
```

```
plt.tight_layout()
plt.show()
```



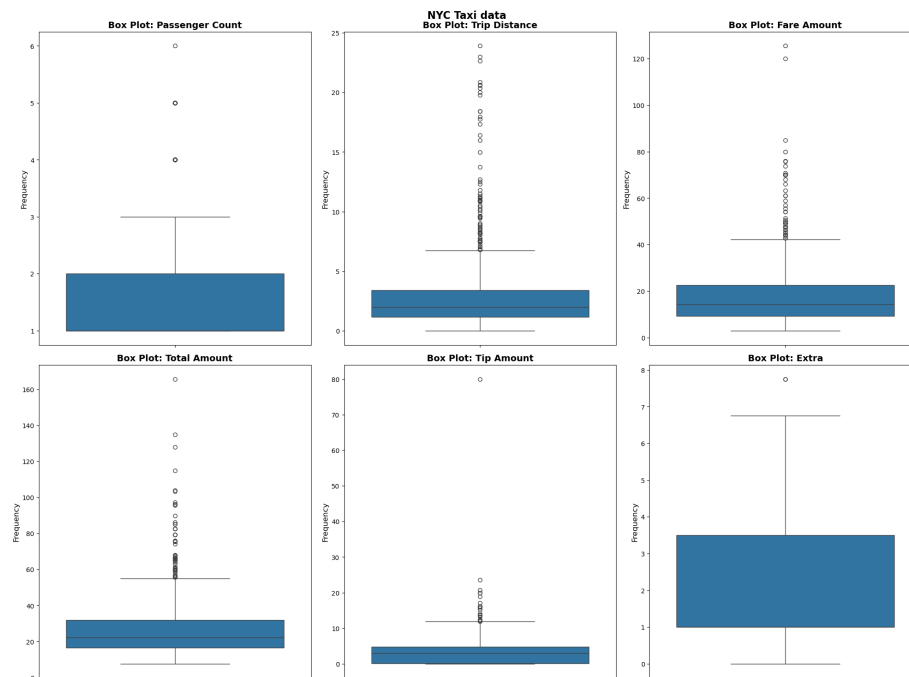
All Observations are same as they were in histogram : 1) Single passenger travels more 2) Most people travel within range of 0-5 km 3) Generally fares are below \$40 4) Total amount comes between \$0 - \$60 5) people generally tips below \$ 10 6) Extras are either 1 or 4

```
fig, axes = plt.subplots(2, 3, figsize = (20,15))
fig.suptitle('NYC Taxi data', fontsize = 16, fontweight = 'bold')

axes = axes.flatten()

for i, col in enumerate(numerical_cols):
    sns.boxplot(y=df_clean[col], ax=axes[i])
    axes[i].set_title(f'Box Plot: {col.replace("_", " ").title()}', fontsize = 14, fontweight = 'bold')
    axes[i].set_ylabel('Frequency', fontsize = 12)

plt.tight_layout()
plt.show()
```



Observations :

1) There are 3 outlier points , Q1 is 1 and Q3 is 2 outside 3 are outliers 2) There are many outliers , Q1 is around 1 median is around 3 Q3 is around 4 and outside 0-7 are outliers 3) There are many outliers , Q1 is around 9 median is around 18 Q3 is around 22 and outside 0-42 are outliers 4) There are many outliers , Q1 is around 16 median is around 27 Q3 is around 31 and outside 10-55 are outliers 5) There are few outliers , Q1 is 0 median is around 3 Q3 is around 5 and outside 12 are outliers 6) There is one outliers , Q1 is 1 median is around 2 Q3 is around 3.5 and outside 7 are outlier

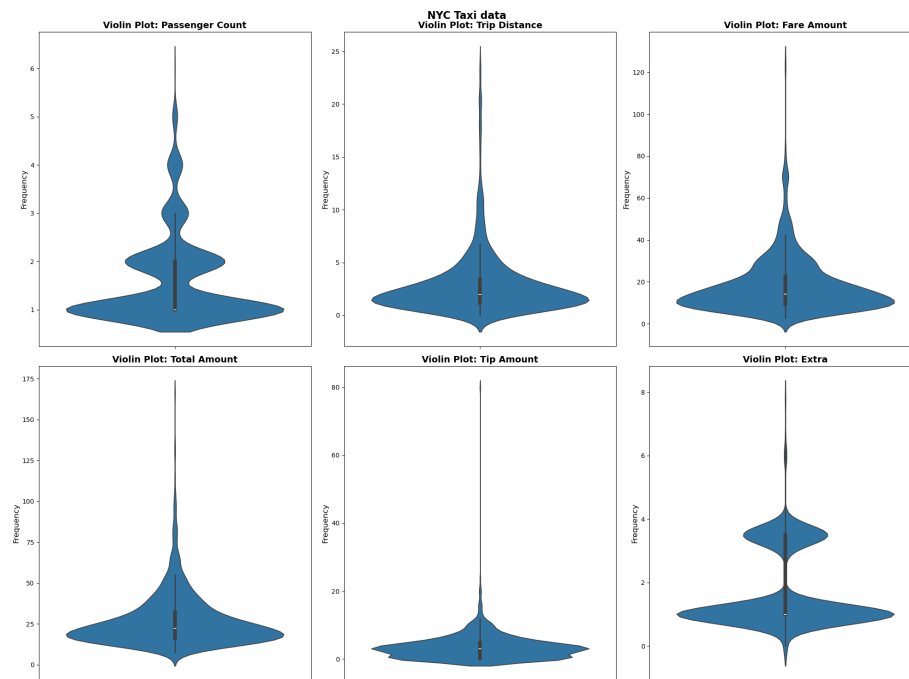
```
fig, axes = plt.subplots(2, 3, figsize = (20,15))
fig.suptitle('NYC Taxi data', fontsize = 16, fontweight = 'bold')
```

```
axes = axes.flatten()
```

```
for i, col in enumerate(numerical_cols):
    sns.violinplot(y=df_clean[col], ax=axes[i])
    axes[i].set_title(f'Violin Plot: {col.replace("_", " ").title()}', fontsize = 14, fontweight = 'bold')
    axes[i].set_ylabel('Frequency', fontsize = 12)
```

```
plt.tight_layout()
plt.show()
```





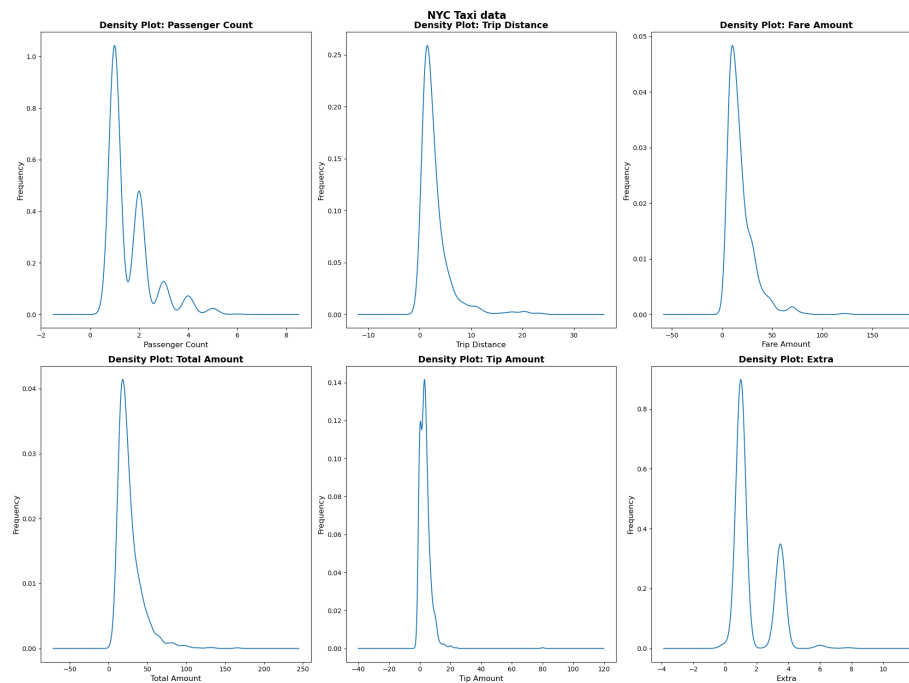
Observation : Mixture of frequency distribution with box plot Wherever the graph is stretch have more frequency 1) Most frequency in about 1 2) Most frequency is about 1 3) Most frequency is about 7 4) Most frequency is about 16 5) Most frequency is about 0 6) Most frequency is about 1

```
fig, axes = plt.subplots(2, 3, figsize = (20,15))
fig.suptitle('NYC Taxi data', fontsize = 16, fontweight = 'bold')
```

```
axes = axes.flatten()
```

```
for i, col in enumerate(numerical_cols):
    df_clean[col].plot(kind='density', ax=axes[i])
    axes[i].set_title(f'Density Plot: {col.replace("_", " ").title()}', fontsize = 14, fontweight = 'bold')
    axes[i].set_xlabel(col.replace("_", " ").title(), fontsize = 12)
    axes[i].set_ylabel('Frequency', fontsize = 12)
```

```
plt.tight_layout()
plt.show()
```



Observation : Shows where the density is highest 1) Most density is about 1 2) Most density is about 1 3) Most density is about 7 4) Most density is about 16 5) Most density is about 0 6) Most density is about 1

*# Create mapping dictionaries for better labels*

```
payment_type_map = {
    1: 'Credit Card',
    2: 'Cash',
    3: 'No Charge',
    4: 'Dispute',
    5: 'Unknown',
    6: 'Voided Trip'
}

ratecode_map = {
    1: 'Standard Rate',
    2: 'JFK',
    3: 'Newark',
    4: 'Nassau/Westchester',
    5: 'Negotiated Fare',
    6: 'Group Ride'
}

vendor_map = {
```

```

        1: 'Creative Mobile Tech',
        2: 'VeriFone Inc'
    }

    flag_map = {
        'Y': 'Yes - Stored',
        'N': 'No - Not Stored'
    }

    fig, axes = plt.subplots(2, 2, figsize=(20, 12))
    fig.suptitle('NYC Taxi Data: Categorical Variables - Bar Charts', fontsize=16, fontweight='bold')

    axes = axes.flatten()

    # Payment Type Distribution
    pymnt_count = df_clean['payment_type'].value_counts()
    payment_labels = [payment_type_map.get(code, f'Code {code}') for code in pymnt_count.index]
    axes[0].bar(payment_labels, pymnt_count.values, color=sns.color_palette("Set2", len(pymnt_count.index)))
    axes[0].set_title('Payment Type Distribution', fontsize=14, fontweight='bold')
    axes[0].set_xlabel('Payment Type', fontsize=12)
    axes[0].set_ylabel('Count', fontsize=12)
    axes[0].grid(True, alpha=0.3)
    axes[0].tick_params(axis='x', rotation=45)

    for i, v in enumerate(pymnt_count.values):
        axes[0].text(i, v + 50, str(v), ha='center', va='bottom')

    # Rate Code ID Distribution
    rcid_count = df_clean['RatecodeID'].value_counts()
    ratecode_labels = [ratecode_map.get(code, f'Code {code}') for code in rcid_count.index]
    axes[1].bar(ratecode_labels, rcid_count.values, color=sns.color_palette("Set2", len(rcid_count.index)))
    axes[1].set_title('Rate Code Distribution', fontsize=14, fontweight='bold')
    axes[1].set_xlabel('Rate Code Type', fontsize=12)
    axes[1].set_ylabel('Count', fontsize=12)
    axes[1].grid(True, alpha=0.3)
    axes[1].tick_params(axis='x', rotation=45)

    for i, v in enumerate(rcid_count.values):
        axes[1].text(i, v + 50, str(v), ha='center', va='bottom')

    # Vendor ID Distribution
    vend_count = df_clean['VendorID'].value_counts()
    vendor_labels = [vendor_map.get(code, f'Vendor {code}') for code in vend_count.index]
    axes[2].bar(vendor_labels, vend_count.values, color=sns.color_palette("Set2", len(vend_count.index)))
    axes[2].set_title('Vendor Distribution', fontsize=14, fontweight='bold')
    axes[2].set_xlabel('Vendor Name', fontsize=12)

```

```

axes[2].set_ylabel('Count', fontsize=12)
axes[2].grid(True, alpha=0.3)
axes[2].tick_params(axis='x', rotation=45)

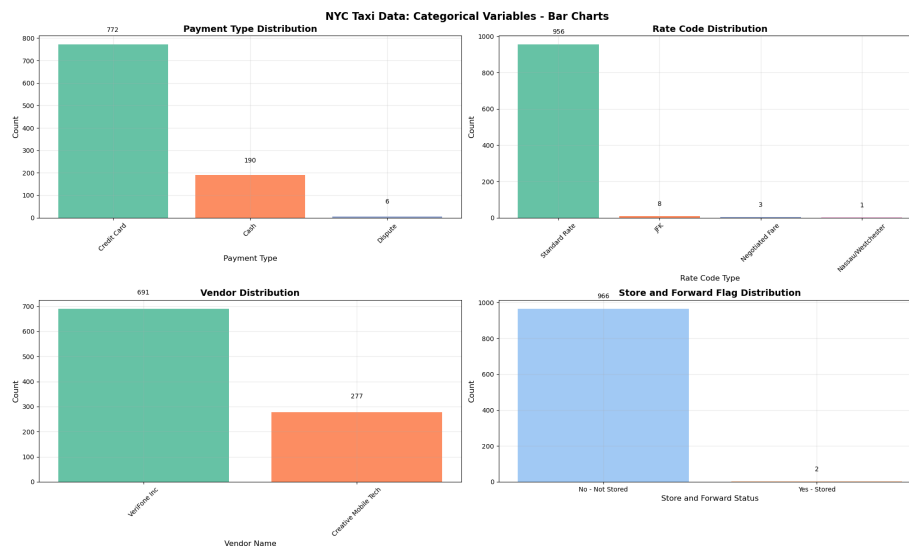
for i, v in enumerate(vend_count.values):
    axes[2].text(i, v + 50, str(v), ha='center', va='bottom')

# Store and Forward Flag Distribution
flag_counts = df_clean['store_and_fwd_flag'].value_counts()
flag_labels = [flag_map.get(flag, flag) for flag in flag_counts.index]
axes[3].bar(flag_labels, flag_counts.values,
            color=sns.color_palette("pastel", len(flag_counts)))
axes[3].set_title('Store and Forward Flag Distribution', fontsize=14, fontweight='bold')
axes[3].set_xlabel('Store and Forward Status', fontsize=12)
axes[3].set_ylabel('Count', fontsize=12)
axes[3].grid(True, alpha=0.3)

for i, v in enumerate(flag_counts.values):
    axes[3].text(i, v + 50, str(v), ha='center', va='bottom')

plt.tight_layout()
plt.show()

```



Observations:

#### 1. PAYMENT TYPE:

Credit Card: 772 trips (79.8%) Cash: 190 trips (19.6%) Dispute: 6 trips (0.6%)

#### 1. RATE CODE ID:

Standard Rate: 956 trips (98.8%) JFK: 8 trips (0.8%) Negotiated Fare: 3 trips (0.3%) Nassau/Westchester: 1 trips (0.1%)

1. VENDOR ID:

VeriFone Inc: 691 trips (71.4%) Creative Mobile Tech: 277 trips (28.6%)

1. STORE AND FORWARD FLAG:

No - Not Stored: 966 trips (99.8%) Yes - Stored: 2 trips (0.2%)

```
fig, axes = plt.subplots(2, 2, figsize=(20, 12))
fig.suptitle('NYC Taxi Data: Categorical Variables - Pie Charts', fontsize=16, fontweight='b')

# Payment Type Pie Chart
payment_counts = df_clean['payment_type'].value_counts()
payment_labels = [payment_type_map.get(x, f'Code {x}') for x in payment_counts.index]
colors1 = sns.color_palette("Set2", len(payment_counts))

axes[0,0].pie(payment_counts.values, labels=payment_labels, autopct='%1.1f%%',
              colors=colors1, startangle=90)
axes[0,0].set_title('Payment Type Distribution', fontsize=14, fontweight='bold')

# Vendor ID Pie Chart
vendor_counts = df_clean['VendorID'].value_counts()
vendor_labels = [vendor_map.get(x, f'Vendor {x}') for x in vendor_counts.index]
colors2 = sns.color_palette("Set3", len(vendor_counts))

axes[0,1].pie(vendor_counts.values, labels=vendor_labels, autopct='%1.1f%%',
              colors=colors2, startangle=90)
axes[0,1].set_title('Vendor Distribution', fontsize=14, fontweight='bold')

# Rate Code ID Pie Chart
ratecode_counts = df_clean['RatecodeID'].value_counts()
ratecode_labels = [ratecode_map.get(x, f'Code {x}') for x in ratecode_counts.index]
colors3 = sns.color_palette("viridis", len(ratecode_counts))

axes[1,0].pie(ratecode_counts.values, labels=ratecode_labels, autopct='%1.1f%%',
              colors=colors3, startangle=90)
axes[1,0].set_title('Rate Code Distribution', fontsize=14, fontweight='bold')

# Store and Forward Flag Pie Chart
if 'store_and_fwd_flag' in df_clean.columns:
    flag_counts = df_clean['store_and_fwd_flag'].value_counts()
    flag_labels = [flag_map.get(x, x) for x in flag_counts.index]
    colors4 = sns.color_palette("pastel", len(flag_counts))

    axes[1,1].pie(flag_counts.values, labels=flag_labels, autopct='%1.1f%%',
```


```

        colors=colors4, startangle=90)
    axes[1,1].set_title('Store and Forward Flag Distribution', fontsize=14, fontweight='bold')
else:
    # Show top 5 pickup locations with zone names
    pickup_top5 = df_clean['PULocationID'].value_counts().head(5)
    pickup_labels = [f'Zone {x}' for x in pickup_top5.index]
    colors4 = sns.color_palette("muted", len(pickup_top5))

    axes[1,1].pie(pickup_top5.values, labels=pickup_labels, autopct='%1.1f%%',
                  colors=colors4, startangle=90)
    axes[1,1].set_title('Top 5 Pickup Locations', fontsize=14, fontweight='bold')

plt.tight_layout()
plt.show()

```



#### 1. PAYMENT TYPE:

Credit Card: 772 trips (79.8%) Cash: 190 trips (19.6%) Dispute: 6 trips (0.6%)

#### 1. RATE CODE ID:

Standard Rate: 956 trips (98.8%) JFK: 8 trips (0.8%) Negotiated Fare: 3 trips (0.3%) Nassau/Westchester: 1 trips (0.1%)

1. VENDOR ID:

VeriFone Inc: 691 trips (71.4%) Creative Mobile Tech: 277 trips (28.6%)

1. STORE AND FORWARD FLAG:

No - Not Stored: 966 trips (99.8%) Yes - Stored: 2 trips (0.2%)

## B - Inferential

```
from scipy import stats
```

```
def cnfdnc_intv(data, confidence = 0.95):
    n = len(data)
    mean = np.mean(data)
    std_err = stats.sem(data)
    h = std_err * stats.t.ppf((1 + confidence) / 2, n - 1)
    return mean - h, mean + h, mean, std_err
```

*# Standard error of the mean*  
*# t-distribution*  
*# CI =  $x \pm (t_{\{ / 2, df \}} \times SE)$*

```
variables = ['trip_distance', 'fare_amount', 'tip_amount']
ci_results = {}
for var in variables:
    data = df_clean[var].dropna()
    lower, upper, mean_val, std_err = cnfdnc_intv(data)
    ci_results[var] = {
        'mean': mean_val,
        'lower_bound': lower,
        'upper_bound': upper,
        'std_error': std_err,
        'sample_size': len(data)
    }
    print(f"\n{var.replace('_', ' ').title()}:")
    print(f"  Sample Size: {len(data)}")
    print(f"  Sample Mean: ${mean_val:.4f}")
    print(f"  Standard Error: ${std_err:.4f}")
    print(f"  95% CI: [${lower:.4f}, ${upper:.4f}]")
    print(f"  Interpretation: We are 95% confident that the true population mean")
    print(f"  {var.replace('_', ' ')} lies between ${lower:.4f} and ${upper:.4f}")
```

Trip Distance:

Sample Size: 968

Sample Mean: \$2.9154

Standard Error: \$0.1000

95% CI: [\$2.7192, \$3.1115]

Interpretation: We are 95% confident that the true population mean trip distance lies between \$2.7192 and \$3.1115

Fare Amount:

Sample Size: 968

Sample Mean: \$18.4021

Standard Error: \$0.4363

95% CI: [\$17.5459, \$19.2582]

Interpretation: We are 95% confident that the true population mean fare amount lies between \$17.5459 and \$19.2582

Tip Amount:

Sample Size: 968

Sample Mean: \$3.4819

Standard Error: \$0.1311

95% CI: [\$3.2247, \$3.7391]

Interpretation: We are 95% confident that the true population mean tip amount lies between \$3.2247 and \$3.7391

*# Hypothesis Test 1: One-sample t-test for tip amount*

```
tip_data = df_clean['tip_amount'].dropna()
```

```
null_hypothesis_value = 2.0      # H0: = $2, H1:  $2
```

```
t_stat, p_value = stats.ttest_1samp(tip_data, null_hypothesis_value)      #  $t = (x - ) /$ 
```

```
print("HYPOTHESIS TEST 1: One-Sample t-test for Tip Amount")
```

```
print(f"\nH0: = ${null_hypothesis_value} (null hypothesis)")
```

```
print(f"H1:  ${null_hypothesis_value} (alternative hypothesis)")
```

```
print(f"Significance level: = 0.05")
```

```
print("\nResults:")
```

```
print(f"Sample size: {len(tip_data)}")
```

```
print(f"Sample mean: ${tip_data.mean():.4f}")
```

```
print(f"Sample std: ${tip_data.std():.4f}")
```

```
print(f"t-statistic: {t_stat:.4f}")
```

```
print(f"p-value: {p_value:.6f}")
```

```
if p_value < 0.05:
```

```
    print("Decision: Reject H0")
```

```
    print(f"Conclusion: There is significant evidence that the average tip amount is different from $2.0")
```

```
else:
```

```
    print("Decision: Fail to reject H0")
```

```
    print(f"Conclusion: There is insufficient evidence that the average tip amount is different from $2.0")
```

HYPOTHESIS TEST 1: One-Sample t-test for Tip Amount

H0: = \$2.0 (null hypothesis)

H1: \$2.0 (alternative hypothesis)

Significance level: = 0.05



```

Results:
Sample size: 968
Sample mean: $3.4819
Sample std: $4.0774
t-statistic: 11.3077
p-value: 0.000000
Decision: Reject H0
Conclusion: There is significant evidence that the average tip amount is different from $2

# Hypothesis Test 2: Two-sample t-test for fare amount by payment type
# Compare credit card (1) vs cash (2) payments

credit_fares = df_clean[df_clean['payment_type'] == 1]['fare_amount'].dropna()
cash_fares = df_clean[df_clean['payment_type'] == 2]['fare_amount'].dropna()

# Perform two-sample t-test
#  $(x_1 - x_2) - (u_1 - u_2)$ 
# -----
# ----- Underroot
#  $/ (s_1^2 / n_1) + (s_2^2 / n_2)$ 
t_stat2, p_value2 = stats.ttest_ind(credit_fares, cash_fares)

print("H0: _credit = _cash (no difference in mean fare amounts)")
print("H1: _credit _cash (difference in mean fare amounts)")
print(f"Significance level: = 0.05")
print()
print("Results:")
print(f"Credit card payments:")
print(f"  Sample size: {len(credit_fares)}")
print(f"  Sample mean: ${credit_fares.mean():.4f}")
print(f"  Sample std: ${credit_fares.std():.4f}")
print()
print(f"Cash payments:")
print(f"  Sample size: {len(cash_fares)}")
print(f"  Sample mean: ${cash_fares.mean():.4f}")
print(f"  Sample std: ${cash_fares.std():.4f}")
print()
print(f"t-statistic: {t_stat2:.4f}")
print(f"p-value: {p_value2:.6f}")
print()

if p_value2 < 0.05:
    print("Decision: Reject H0")
    print("Conclusion: There is significant evidence of a difference in mean fare amounts be
else:

```

```

    print("Decision: Fail to reject H0")
    print("Conclusion: There is insufficient evidence of a difference in mean fare amounts b

H0: _credit = _cash (no difference in mean fare amounts)
H1: _credit _cash (difference in mean fare amounts)
Significance level: = 0.05

Results:
Credit card payments:
    Sample size: 772
    Sample mean: $18.2481
    Sample std: $13.1342

Cash payments:
    Sample size: 190
    Sample mean: $18.9542
    Sample std: $15.1311

t-statistic: -0.6435
p-value: 0.520065

Decision: Fail to reject H0
Conclusion: There is insufficient evidence of a difference in mean fare amounts between paym

from scipy.stats import chi2_contingency

# Hypothesis Test 3: Chi-square test of independence
# Test if Payment type and RateCodeID are independent

# Create contingency table
contingency_table = pd.crosstab(df_clean['payment_type'], df_clean['RatecodeID'])
print("HYPOTHESIS TEST 3: Chi-square Test of Independence")
print("\nH0: Payment type and Rate code are independent")
print("H1: Payment type and Rate code are not independent")
print(f"Significance level: = 0.05")
print("Contingency Table:")
print(contingency_table)

# Perform chi-square test
#  $\chi^2 = [(Observed - Expected)^2 / Expected]$ 
chi2_stat, p_value3, dof, expected = chi2_contingency(contingency_table)

print("Results:")
print(f"Chi-square statistic: {chi2_stat:.4f}")
print(f"Degrees of freedom: {dof}")
print(f"p-value: {p_value3:.6f}")

```

```

if p_value3 < 0.05:
    print("Decision: Reject H0")
    print("Conclusion: Payment type and Rate code are NOT independent")
else:
    print("Decision: Fail to reject H0")
    print("Conclusion: Payment type and Rate code are independent")

```

#### HYPOTHESIS TEST 3: Chi-square Test of Independence

H0: Payment type and Rate code are independent  
H1: Payment type and Rate code are not independent  
Significance level: = 0.05

Contingency Table:

RatecodeID	1.0	2.0	4.0	5.0
payment_type				
1	765	6	0	1
2	185	2	1	2
4	6	0	0	0

Results:

Chi-square statistic: 8.5424

Degrees of freedom: 6

p-value: 0.200995

Decision: Fail to reject H0

Conclusion: Payment type and Rate code are independent

```

print("Correlation:")
print("Pearson - linear data")
print(f"Trip Distance vs Fare Amount: {df_clean['trip_distance'].corr(df_clean['fare_amount'])}")
print(f"Fare Amount vs Tip Amount: {df_clean['fare_amount'].corr(df_clean['tip_amount'])}")
print("spearman - non linera data")
print(f"Fare Amount vs Tip Amount: {df_clean['fare_amount'].corr(df_clean['tip_amount'], method='spearman')}")

```

```

def interpret_correlation(r):
    if abs(r) >= 0.8:
        return "very strong"
    elif abs(r) >= 0.6:
        return "strong"
    elif abs(r) >= 0.4:
        return "moderate"
    elif abs(r) >= 0.2:
        return "weak"
    else:
        return "very weak"

```

```

print("Interpretation")
print(f"- Trip distance and fare amount show a {interpret_correlation(df_clean['trip_distance'].corr(df_clean['fare_amount']))} correlation")
print(f"- Fare amount and tip amount show a {interpret_correlation(df_clean['fare_amount'].corr(df_clean['tip_amount']))} correlation")

```

```

Correlation:
Pearson - linear data
Trip Distance vs Fare Amount: 0.9045
Fare Amount vs Tip Amount: 0.4442
spearman - non linera data
Fare Amount vs Tip Amount: 0.4099
Interpretation
- Trip distance and fare amount show a very strong positive correlation
- Fare amount and tip amount show a moderate positive correlation

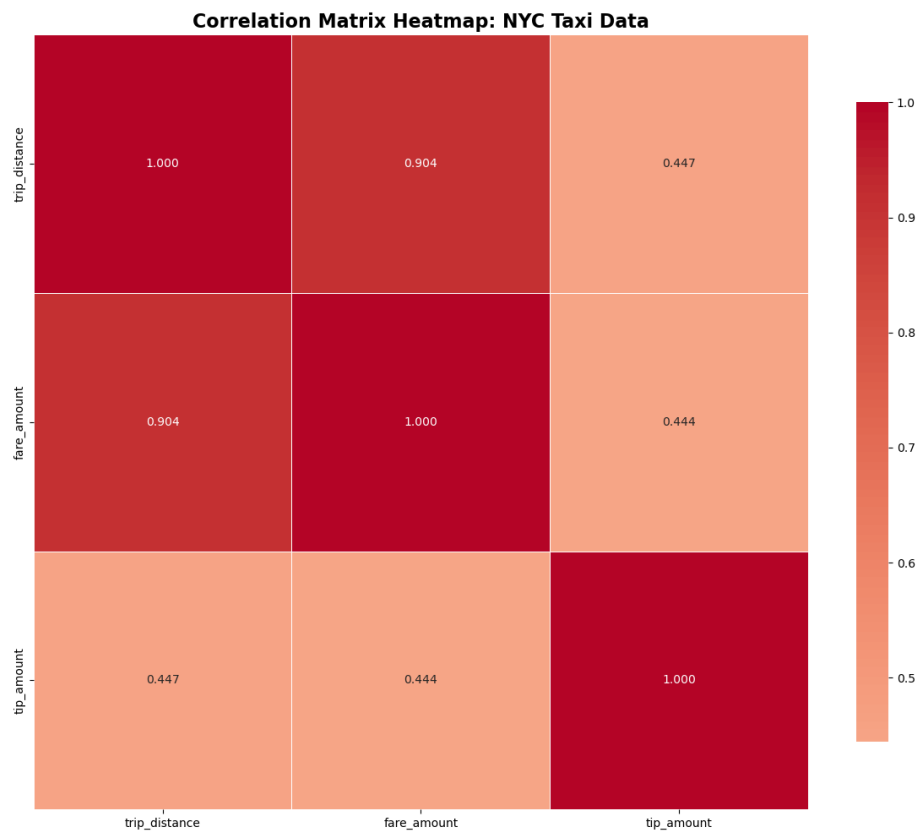
correlation_variables = ['trip_distance', 'fare_amount', 'tip_amount']
correlation_matrix = df_clean[correlation_variables].corr()

plt.figure(figsize=(12, 10))

# Create heatmap
sns.heatmap(correlation_matrix,
            annot=True,
            cmap='coolwarm',
            center=0,
            square=True,
            linewidths=0.5,
            cbar_kws={"shrink": 0.8},
            fmt='.3f')

plt.title('Correlation Matrix Heatmap: NYC Taxi Data', fontsize=16, fontweight='bold')
plt.tight_layout()
plt.show()

```



## Bonus

```
df_clean['tpep_pickup_datetime'] = pd.to_datetime(df_clean['tpep_pickup_datetime'])
df_clean['tpep_dropoff_datetime'] = pd.to_datetime(df_clean['tpep_dropoff_datetime'])

# Extract hour of day for pickup
df_clean['pickup_hour'] = df_clean['tpep_pickup_datetime'].dt.hour

# Trip count by hour
hourly_trips = df_clean['pickup_hour'].value_counts().sort_index()

# Fare amount by hour
hourly_fare = df_clean.groupby('pickup_hour')['fare_amount'].mean()

# Create time series plots
fig, axes = plt.subplots(2, 2, figsize=(16, 10))

# Plot 1: Trip count by hour
```

```

axes[0,0].bar(hourly_trips.index, hourly_trips.values, color='skyblue')
axes[0,0].set_title('Trip Count by Hour of Day')
axes[0,0].set_xlabel('Hour of Day')
axes[0,0].set_ylabel('Number of Trips')
axes[0,0].grid(True, alpha=0.3)

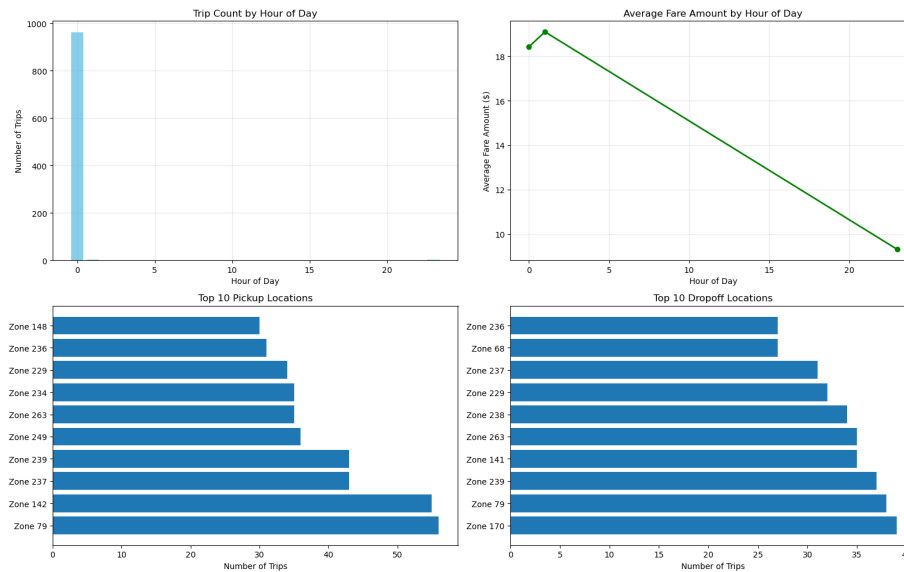
# Plot 2: Average fare by hour
axes[0,1].plot(hourly_fare.index, hourly_fare.values, marker='o', linewidth=2, color='green')
axes[0,1].set_title('Average Fare Amount by Hour of Day')
axes[0,1].set_xlabel('Hour of Day')
axes[0,1].set_ylabel('Average Fare Amount ($)')
axes[0,1].grid(True, alpha=0.3)

# Plot 3: Pickup location distribution
pickup_zones = df_clean['PULocationID'].value_counts().head(10)
axes[1,0].barh(range(len(pickup_zones)), pickup_zones.values)
axes[1,0].set_yticks(range(len(pickup_zones)))
axes[1,0].set_yticklabels([f'Zone {x}' for x in pickup_zones.index])
axes[1,0].set_title('Top 10 Pickup Locations')
axes[1,0].set_xlabel('Number of Trips')

# Plot 4: Dropoff location distribution
dropoff_zones = df_clean['DOLocationID'].value_counts().head(10)
axes[1,1].barh(range(len(dropoff_zones)), dropoff_zones.values)
axes[1,1].set_yticks(range(len(dropoff_zones)))
axes[1,1].set_yticklabels([f'Zone {x}' for x in dropoff_zones.index])
axes[1,1].set_title('Top 10 Dropoff Locations')
axes[1,1].set_xlabel('Number of Trips')

plt.tight_layout()
plt.show()

```



## Conclusion

### 1. DESCRIPTIVE STATISTICS INSIGHTS:

- Trip Distance: Right-skewed distribution with mean 2.9 miles
- Fare Amount: Positive skew, average around \$17-18
- Tip Amount: Highly right-skewed, many zero values
- Passenger Count: Most trips have 1-2 passengers
- Total Amount: Strong correlation with fare amount

### 1. HYPOTHESIS TESTING RESULTS:

- Average tip amount is significantly different from \$2
- no significant difference in fare amounts between payment types
- Payment type and rate code are independent

### 1. CORRELATION FINDINGS:

- Strong positive correlation between trip distance and fare amount ( $r = 0.904$ )
- Moderate positive correlation between fare amount and tip amount ( $r = 0.444$ )
- Fare amount is the primary component of total amount

### 1. TIME-BASED PATTERNS:

- Peak trip hour: 0:00 with 962 trips
- Lowest trip hour: 1:00 with 3 trips
- Rush hour patterns visible in trip frequency